

Design and Implementation of Database System Internals

①

Welcome & Setup

Winter 2025/26

Torsten Grust
Universität Tübingen, Germany

Dissecting the Duck's Innards

①

Welcome & Setup

Winter 2025/26

Torsten Grust
Universität Tübingen, Germany

1 | Welcome (Back)!



Welcome to this course which is all about digging deep into the internals of **tabular database management systems** (DBMSs).

Our tour through the DBMS kernel will touch on

- the efficient representation of data on **secondary** (SSD/HDD 🗄️) and **primary storage** (RAM 🖨️),
- turning declarative SQL queries into efficient **flows of data**,
- a variety of interesting **data structures** 🗂️ for **sizable volumes** and associated **algorithms**,
- **modern CPUs** 🖨️ and how looping/branching code ↻↪ executes,
- various forms of **parallelism** ≡ on different levels (from single CPU instructions to threads), or
- ensuring **data integrity** under concurrent access or even if the host machine fails 🗨️ 🗑️.


Dissecting the Duck's Innards










This semester we will focus on **DuckDB** , a contemporary tabular DBMS built for high-speed SQL-based data analytics.

- DuckDB is young, first released in June 2019 (while most DBMSs—like PostgreSQL—originate from the 1980s).
- DuckDB development is moving fast. The system both adopts newest research results and builds on established DB wisdom.
- DuckDB is developed in the open (, MIT license).¹ All code is available for inspection and tinkering.
- DuckDB provides a range of hooks to observe its internals.
- DuckDB comes with a CLI  and programming language APIs.
- DuckDB is easy to install and requires no maintenance.

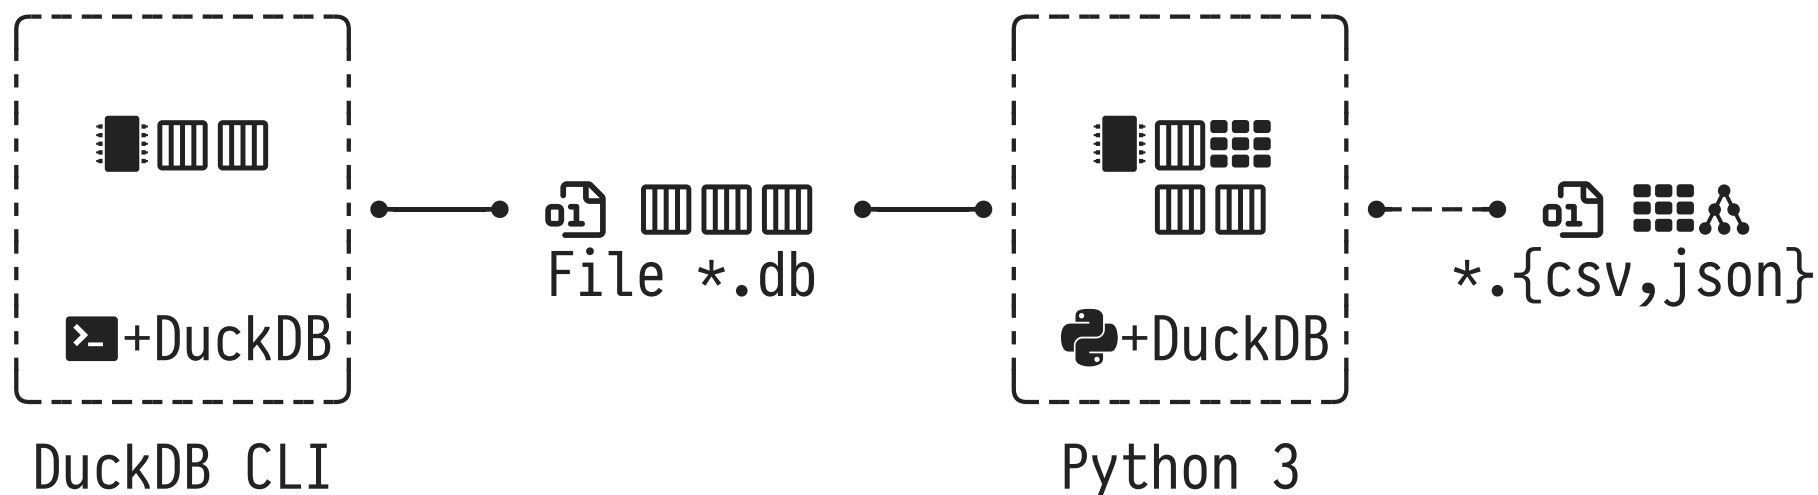
¹ By now, quite a bit of DuckDB's kernel code has been developed here at U Tübingen.

Not DuckDB: Dedicated DBMS Process/Host Isolated From Users



- DBMS  controls family of disks holding table data , maps relevant table fragments into RAM buffer .
 - On-disk data organized in directories  or in raw blocks.
- Client processes  and DBMS process  isolated, connected via network  or inter-process communication .
- Data needs to be de-/serialized after/before wire transfer.
- Data structures in clients  are inaccessible by the DBMS.
- DBMS archetypes: PostgreSQL, MySQL, SQL Server™, Oracle®.

DuckDB: A Tabular DBMS Inside Your Own Process²



- DuckDB kernel and client share a **single process** `[]`.
 - Python: `import duckdb`, C/C++: link with `libduckdb`.
- Table data resides in a **single database file** `o1` *.db, native DuckDB data formats `||||` in file and in RAM are similar.
- DuckDB sees in-process client data and can **directly read/write many client data structures** `|||||` using SQL (“zero copy”).
- DuckDB allocates sizable RAM buffers (but can use disk `📁` for temporary storage if required).


² An frequently used tag line reads: “DuckDB: The SQLite for Analytics.”

2 | This Course (*Dissecting the Duck's Innards*, short: *DiDi*)

- We will focus on DuckDB as a **tabular** SQL-based DBMS 🗃️. There are other kinds of DBMSs (for graphs, key/value pairs, vectors, ...), but we will not discuss those here.
- We will get our hands dirty using **DuckDB** 🐥 and its extensive **SQL** dialect. Lots of SQL will be read and written.
- Whenever possible we try to observe DuckDB under load or use hooks to inspect its operation while SQL queries are processed.
- We thus assume basic familiarity with the tabular data model and SQL, e.g., as discussed in *Tabular Database Systems (TaDa)*.
- We will draw data and queries from a variety of sources and have **fun** along the way! 😊

Torsten Grust?

Time Frame	Affiliation/Position
1989-1994	Diploma in Computer Science, TU Clausthal
1994-1999	Promotion (PhD), U Konstanz
2000	<i>Visiting Researcher</i> , IBM (USA)
2000-2004	Habilitation, U Konstanz
2004-2005	Professor Database Systems, TU Clausthal
2005-2008	Professor Database Systems, TU München
since 2008	Professor Database Systems, U Tübingen

- Web: <https://db.cs.uni-tuebingen.de/grust>
- Office: WSI, Sand 13, Room B318
- Bluesky 🦋: [@tegggy.org](https://bsky.app/profile/tegggy.org)
- Best bet is to catch me on the [DB group's Discord server](#) 

Administrivia





Weekday/Time	Slot	Room
Monday, 10:15–11:45 ☹	Lecture	Lecture Hall TTR2 (MvL 6)
Thursday, 14:15–15:45 ☹	Tutorial	Lecture Hall TTR2 (MvL 6)

- **⚠ No lecture/tutorials on**
 - Thu, October 16
 - (In this winter semester, we may miss Mon, Jan 19, 2026 and Thu, Jan 22, 2026—details will follow in due course.)
- **End-Term Exam (6 ECTS)**
 - Written exam on Mon, Feb 9, 2026, 10:00–12:00 (Room *tbd*)
 - Score $\geq \frac{2}{3}$ of the overall assignment points to be admitted to the exam.



Slides and Further Lecture Material

These **slides** (PDF), **code fragments** (SQL, Python, C), and **sample data** will be uploaded to a GitHub  repository:


github.com/DBatUTuebingen-Teaching/didi-ws2526 

- Slides point to relevant code files or extra material using tags like  #001:
 - Refers to a file named `001-*` on the GitHub repository (e.g., `001-sum-quantity.awk`).
- **NB.** Code and extra material provide essential content (e.g., details on SQL-based experiments or sample data).
 -  +  = : Only slides + code provide a complete picture.

Weekly Assignments & Tutorial Sessions

- We will distribute, collect, and grade **weekly assignments** (Wednesday→Wednesday) via GitHub .
- You work on these in **teams of two**. Hand-in again via GitHub .

Organized and run by **Denis Hirn**:

- Web: <https://db.cs.uni-tuebingen.de/team/members/denis-hirn/>
- E-Mail: db-lehre@cs.uni-tuebingen.de
- Office: WSI, Sand 13, Rooms B314
- Find us on the [Discord server](#) 

Assignments start once we have collected the first batch of interesting material, probably on Wednesday, October 22.

Discord

During this winter semester, the **DB Discord** is *the* course hub:

<https://db.cs.uni-tuebingen.de/discord> 

- ⚠ **Registration** (do it!): use `/verify` with your e-mail address
- ❓ Questions and answers (do *no* post complete solutions)
- ⬇ Download additional code examples (e.g., SQL snippets)
- 💬 General discussion
- ⌚ Quick turnaround (responses often within minutes)

DiDi's Course Homepage

<https://db.cs.uni-tuebingen.de/teaching/ws2526/didi/> 

- **Organizational matters**



Curriculum. General announcements regarding the lecture, exams, rooms, or dates. (Less important this semester.)

- **Contact information**

Turn to Discord first. But feel free to send e-mail if you seek specific help/need to discuss personal issues with us.

Material


This course is *not* based on a single textbook. Rather, we build on

- a variety of scientific papers,
- textbook excerpts (few),
- the DuckDB  documentation at <https://duckdb.org/docs/>,
- Python/C/C++ code snippets (our own and from inside the ),
- blog posts from a range of authors,
- SQL references/standards,
- experience, and best practices.

There is a plethora of books on tabular DBMSs (both usage and internals), sample SQL snippets (experiments, benchmarks, and idioms), or performance tweaks. If we will use such sources, we will provide pointers.

Get Your Hands Dirty: Install DuckDB!

The tabular DBMS **DuckDB** will be the primary tool in this course:



DuckDB

<https://duckdb.org>, version 1.4 (October 2025: 1.4.1)

- Implements an extensive SQL dialect, is highly performant, open to contributions, and generally awesome.
- Straightforward to install and use on macOS 🍏, Windows 🪟, Linux 🐧 (x86 + ARM).

No DuckDB CLI (📄 SQL prompt/REPL) on iOS or Android.³

³ Run the DuckDB CLI in the web browser: <https://shell.duckdb.org>. Suffices for quick SQL experiments.

3 | DuckDB? 🦆?

In case you were wondering:

DuckDB has been named after *Wilbur, the Duck*, which has been living as a pet with Hannes Mühleisen⁴—co-inventor of DuckDB with Mark Raasveldt—on Hannes' houseboat in Amsterdam.

Hannes (CEO) and Mark (CTO) run **DuckDB Labs**, a company that provides support and consultancy services around DuckDB. The labs are located in Amsterdam, The Netherlands.



Hannes and Wilbur (© Hannes Mühleisen)

⁴ Hannes originally is from the Stuttgart area. Back then he used the license plate **S QL 1337**.