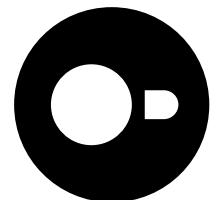


Dissecting the Duck's Innards



④

Sorting Large Tables

Winter 2025/26

Torsten Grust
Universität Tübingen, Germany

1 | Sorting is a Core Operation

The tabular model is based on unordered bags of rows, but the **sorting** of (intermediate) results of SQL queries is material:

1. **ORDER BY** clause,¹
2. ordered aggregates (`list(e1 ORDER BY e2)`),
3. window functions (`row_number() OVER (ORDER BY e)`),
4. joins over time series data (**ASOF JOIN**), or
5. to ensure deterministic **LIMIT** and **OFFSET** behavior.

```
SELECT *          -- wide payload (here: reorder all columns)
FROM lineitem
ORDER BY l_shipdate DESC NULLS FIRST , l_quantity;
--           ↑
-- defaults: ASC NULLS LAST lexicographic ordering
```

¹ In DuckDB, can adapt ordering defaults via `SET default_order = {ASC,DESC}` and `SET default_null_order = {NULLS_FIRST,NULLS_LAST,NULLS_FIRST_ON_ASC_LAST_ON_DESC,NULLS_LAST_ON_ASC_FIRST_ON_DESC}`.

Sorting (Large) Tables in DuckDB

To this day,² DuckDB's internal sorting routines are on the workbench ↗ and still undergo active development.

Key aspects:

- **Efficient row comparison** that respects column types, directions ($\uparrow\downarrow$), handles **NULLs**, allows wide keys (lexicographic ordering).
- Adaptation to **pre-sorted input data**.
- Sorting strategy can process **tables larger than main memory**, spilling to disk 📁 if required.
 - Effectively, sorting needs to materialize its entire input table (last input row could be the first in sort order).
- Balanced **use of all available CPU cores** ⚙ during sorting.

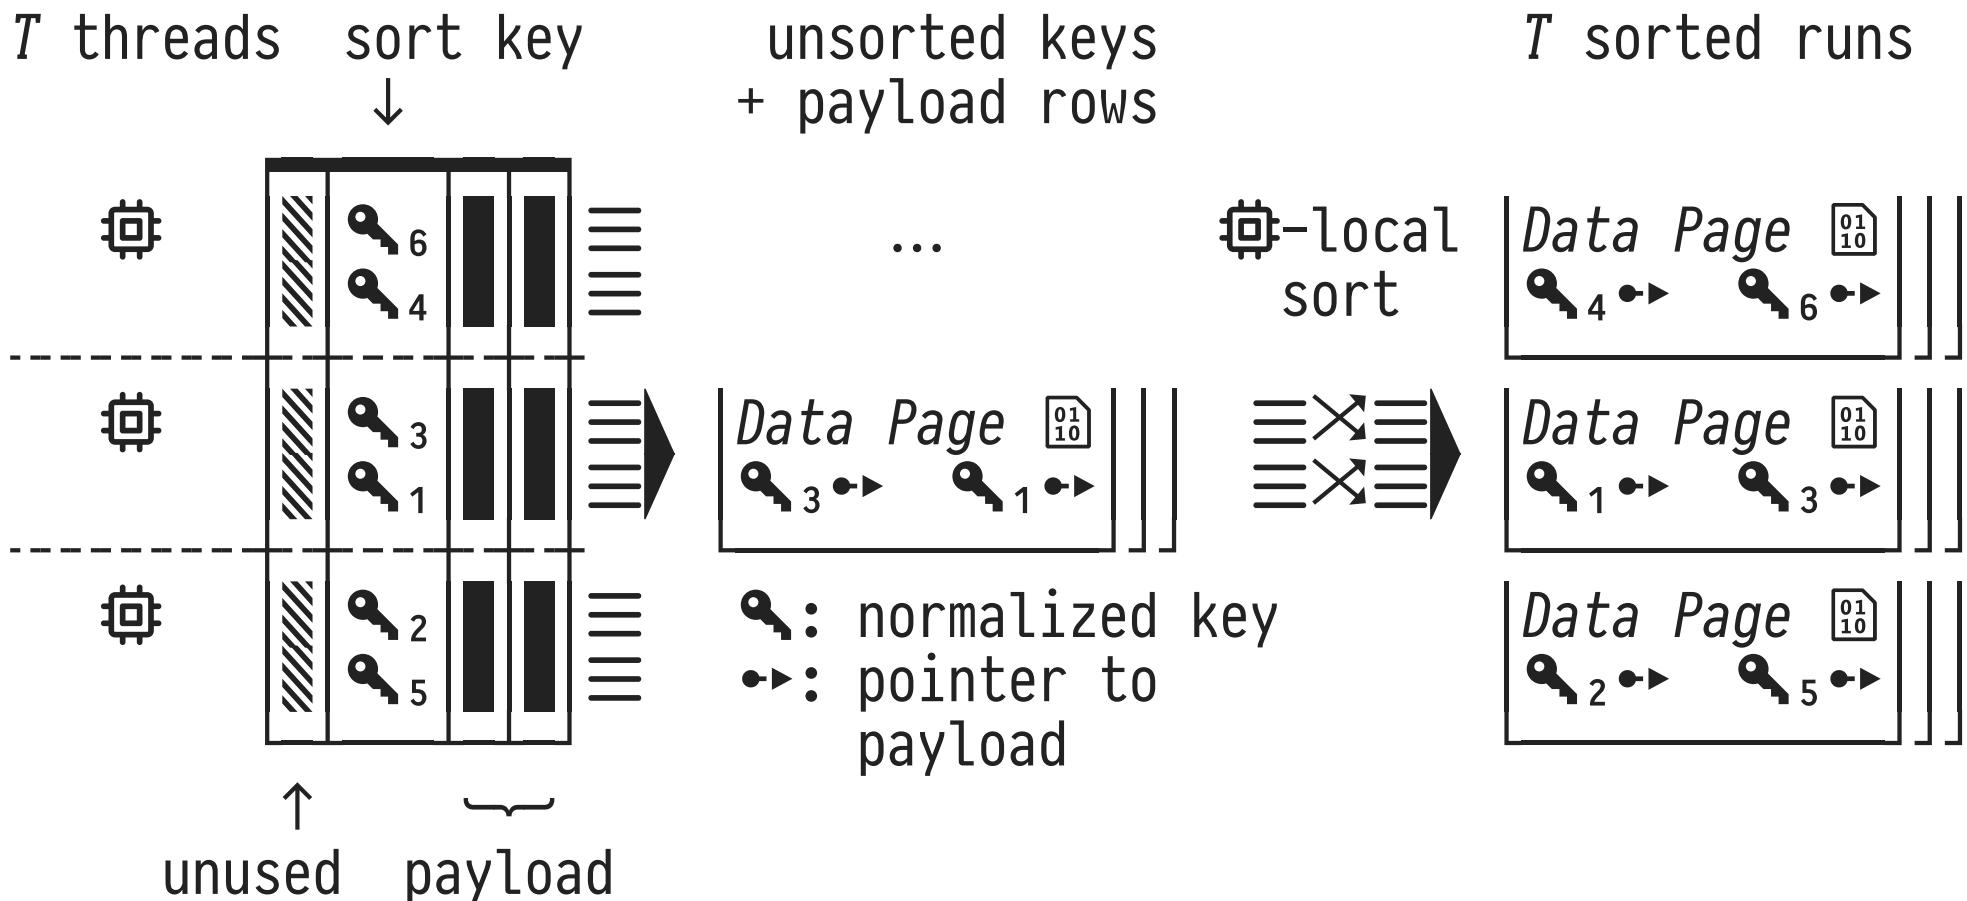
² The last major rewrite of DuckDB's sorting code has shipped with version 1.4.0 in September 2025.

2 | DuckDB: Two-Phase Merge Sort

DuckDB implements a variant of a **two-phase merge sort** strategy that utilizes all available CPU cores  (say T) in both phases:

- **Phase ① (thread-local sorting):**
 - Each : read $1/T$ fraction of input, gather relevant columns to form rows on data pages , **normalize sort keys** ,
 - then generate a **sorted run** of keys.
- **Phase ② (T -way merge of the sorted runs):**
 - Precompute boundaries of T non-overlapping **segments** in each run that can be **merged independently** by the s.
 - Each  immediately outputs its merged run segments for consumption by the downstream query plan.

Phase ①: Thread-Local Sorting



Phase ①: 🔑 Key Normalization

- Sorting compares keys to decide row order \rightarrow . These comparison are frequent, but can be complex and costly:

:

ORDER BY c_1 ASC NULLS LAST, c_2 DESC NULLS FIRST

- Dispatch on type of c_i to select proper $<$ operator, respect sort order ($\uparrow\downarrow$), implement lexicographic ordering (column c_1 dominates), proper **NULL** treatment.
- **⌚ Key normalization:** Map³ values in columns c_i to a byte sequence such that a *single* $<$ comparison on two sequences can decide row order.

⌚ #014

³ DuckDB's key normalization is available at the user level in terms of scalar SQL function [create_sort_key\(\)](#).

Phase ①: Key Normalization for Fixed-Size Keys

CPUs compare fixed-size types (e.g., 32/64-bit ints) faster than variable-length byte sequences.

- Ω If sequence length is known to be n bytes, store it in groups of $\lceil n/8 \rceil$ 64-bit ints. Directly compare these integers:⁴

```
struct FixedSortKey {
    uint64_t      part0;    // assume 8 < n ≤ 16: group size = 2
    uint64_t      part1;
    sort_key_ptr_t payload; // (only if query has payload)
};

bool LessThan(const FixedSortKey &k1, const FixedSortKey &k2) {
    return k1.part0 < k2.part0 ||
           (k1.part0 == k2.part0 && k1.part1 < k2.part1);
}
```

⁴ If SQL queries have no payload (`SELECT c1,c2 FROM ... ORDER BY c1,c2`), DuckDB only keeps the normalized keys in fields `part0`, ...: normalization can be inverted before output is generated.

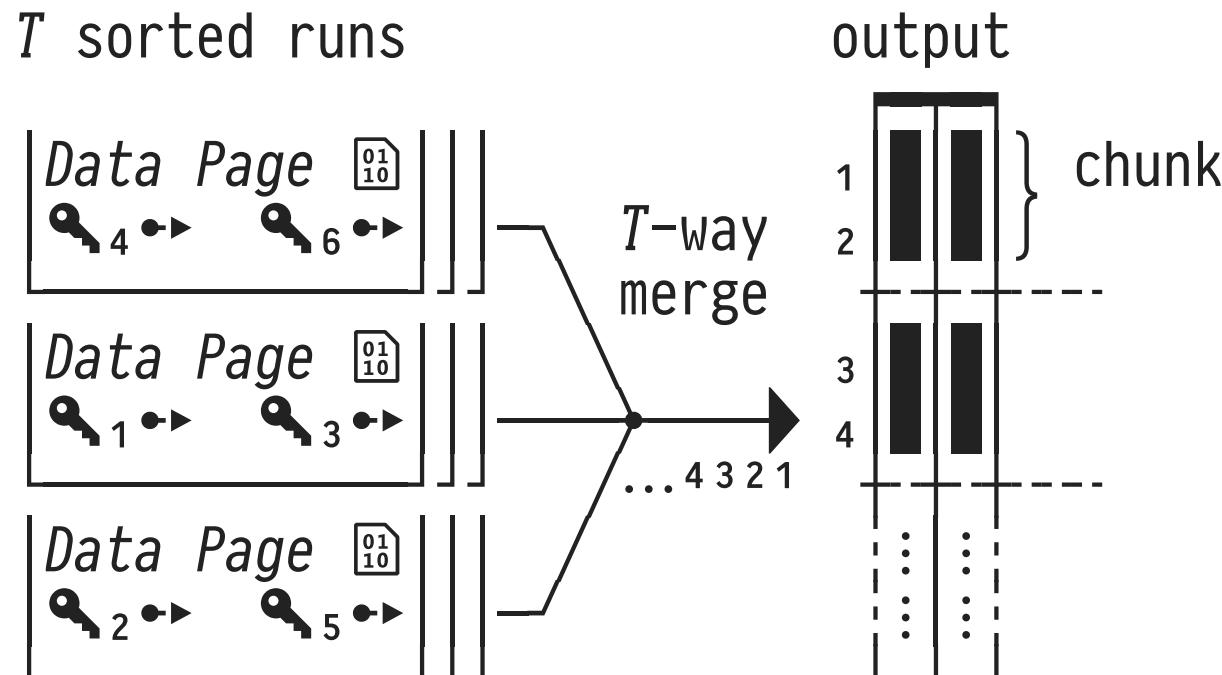
Phase ①: Sorting Algorithms

Typed keys (that can be compared via `<`) fit off-the-shelf sort implementations that assume the C++ `std::iterator` interface.

- DuckDB implements `std::iterator` for chains of 256KB data pages:
 →  → ... → , maps element indices to (page#, offset) pairs.
 - (A single 256KB data page would only be able to hold about 10,000 24-byte `FixedSortKey` structs. DuckDB aims to generate runs much longer than that.)
- DuckDB thus is able to adapt and combine three existing sorting algorithms:
 1. *Vergesort* (detects runs of already sorted data).
 2. *Ska Sort* (radix sort on the first 64 bits of the key).
 3. *Pattern-defeating QuickSort* (fallback, if the first 64 bits do not already order the data).

Phase ②: T-Way Merge

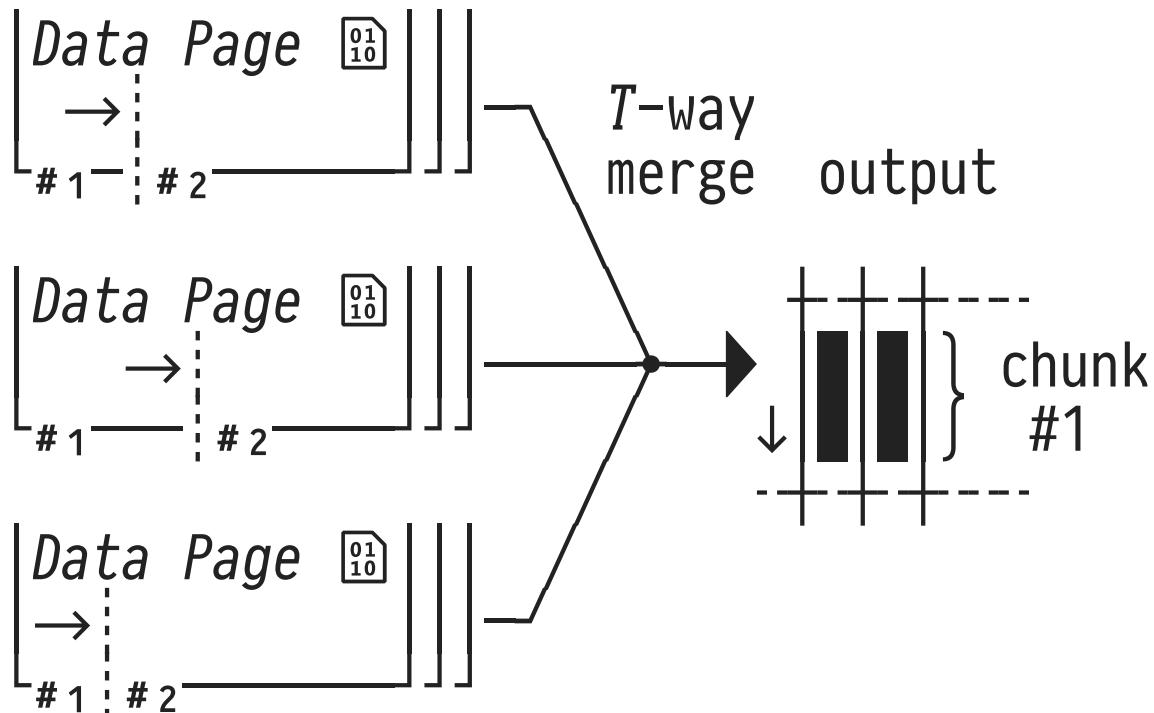
- Consume the T sorted runs produced by Phase ①:
 - **Merge** pairs $(\text{key}_k, \rightarrow)$ using $<$ on key key_k ,
 - dereference pointers \rightarrow to access payload,
 - emit sorted output columns chunk-by-chunk:



- How can all CPU cores  contribute equally during this merge?

Phase ②: Parallelizing T-Way Merge

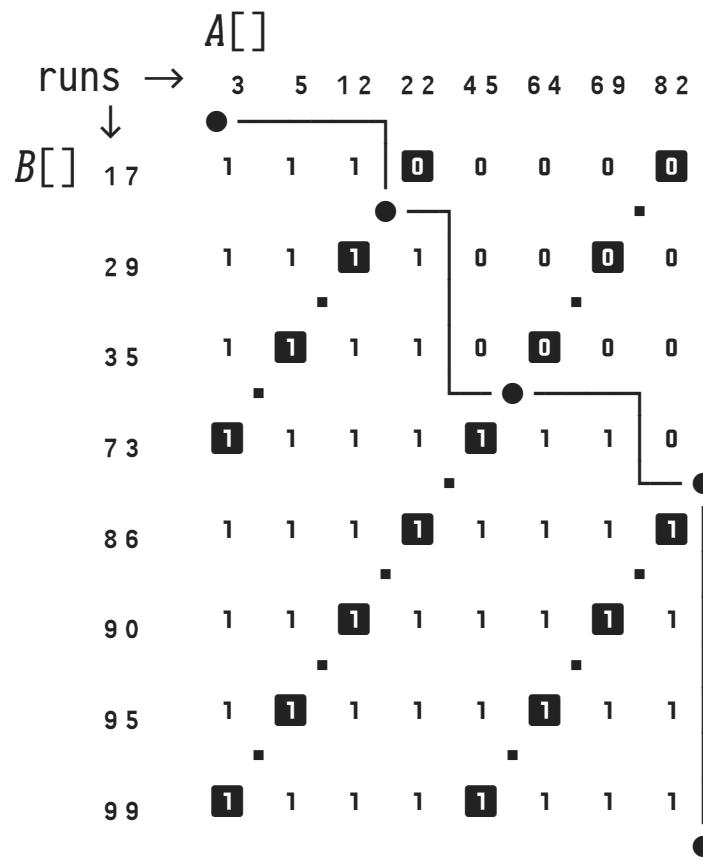
T sorted runs



- Runs contribute rows in **segment** #1 (up to \vdots) to output chunk #1.
- Run segments $\#i$ hold the rows for chunk $\#i$. No overlap between segments.
- Ω Precompute segment boundaries \vdots . Threads \otimes work on one segment.

- Once run segments $\#i$ are merged, the thread \otimes *immediately* emits chunk $\#i$. The downstream plan can assemble the sorted output based on indices $\#i$ (DuckDB lingo: **batch indices**).

Phase ②: Precomputing Segment Boundaries (*Merge Path*)



- Merge matrix $M[i,j]$: 1 if $A[i] < B[j]$, 0 otherwise.
- Merge path
- The n th point of the merge path lies on the n th cross diagonal
- Partitioning the path into k segments of equal size can be done by considering $k-1$ evenly-spaced diagonals.
- Point
- Points

$A:$	3	5	12	22	45	64	69	82	
$B:$	17	29	35	73	86	90	95	99	
	#1	#2	#3		#4				

- DuckDB generalizes the *Merge Path* idea from two runs to T runs. (Profiling shows that boundary computation uses $\leq 2\%$ of the overall execution time.)