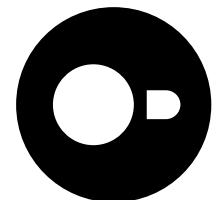


Dissecting the Duck's Innards



⑤

The ART of Indexing

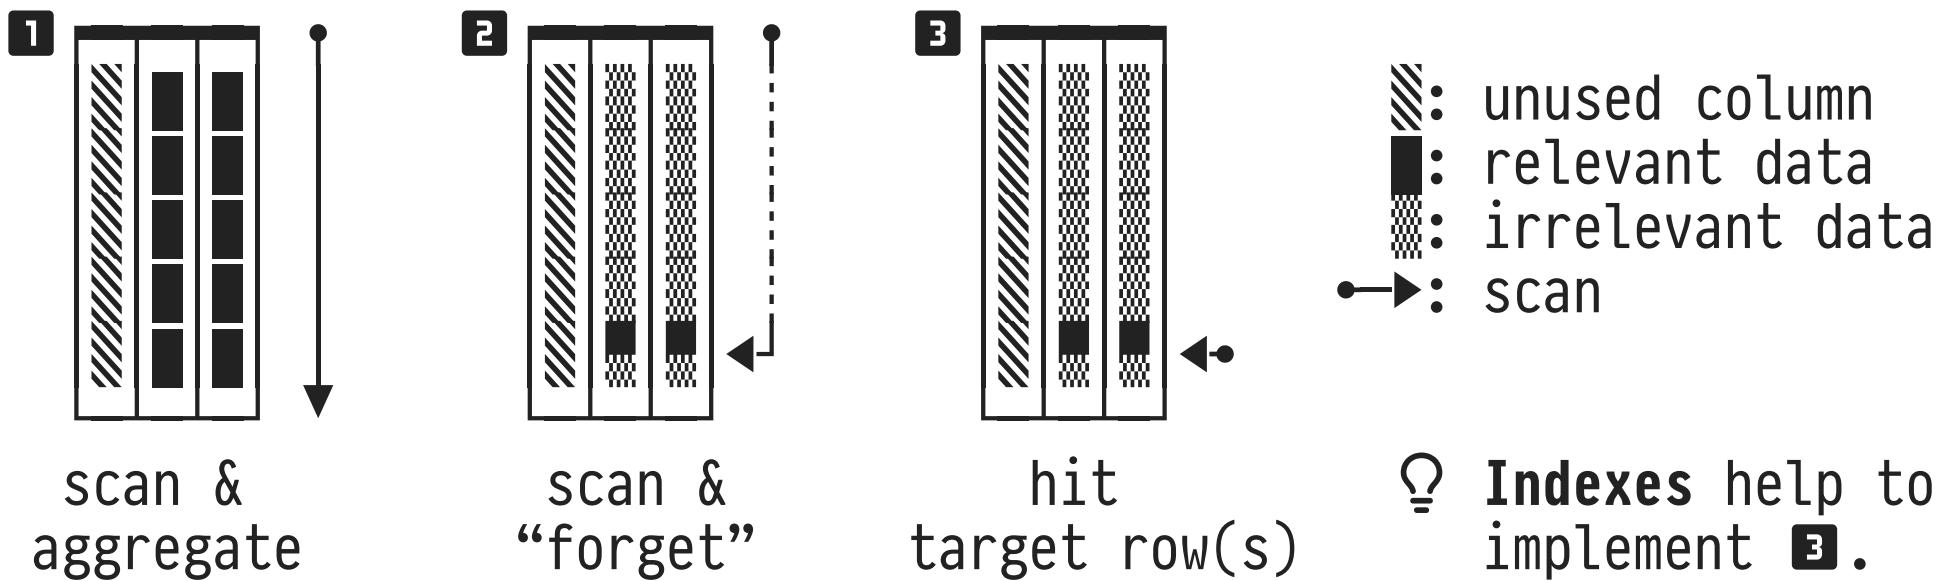
Winter 2025/26

Torsten Grust
Universität Tübingen, Germany

1 | Scanning All Rows vs. Narrowing In On Few Rows

The internals of DuckDB have been designed to efficiently support analytical SQL queries: “*online analytical processing*” (OLAP).

- Typically, these queries read (and aggregate) *all rows* of the input tables 1.
- Yet, scanning all rows wastes I/O and memory bandwidth 2 if a query focuses on *small row subsets* (or even a *single row*):



Indexes in DuckDB

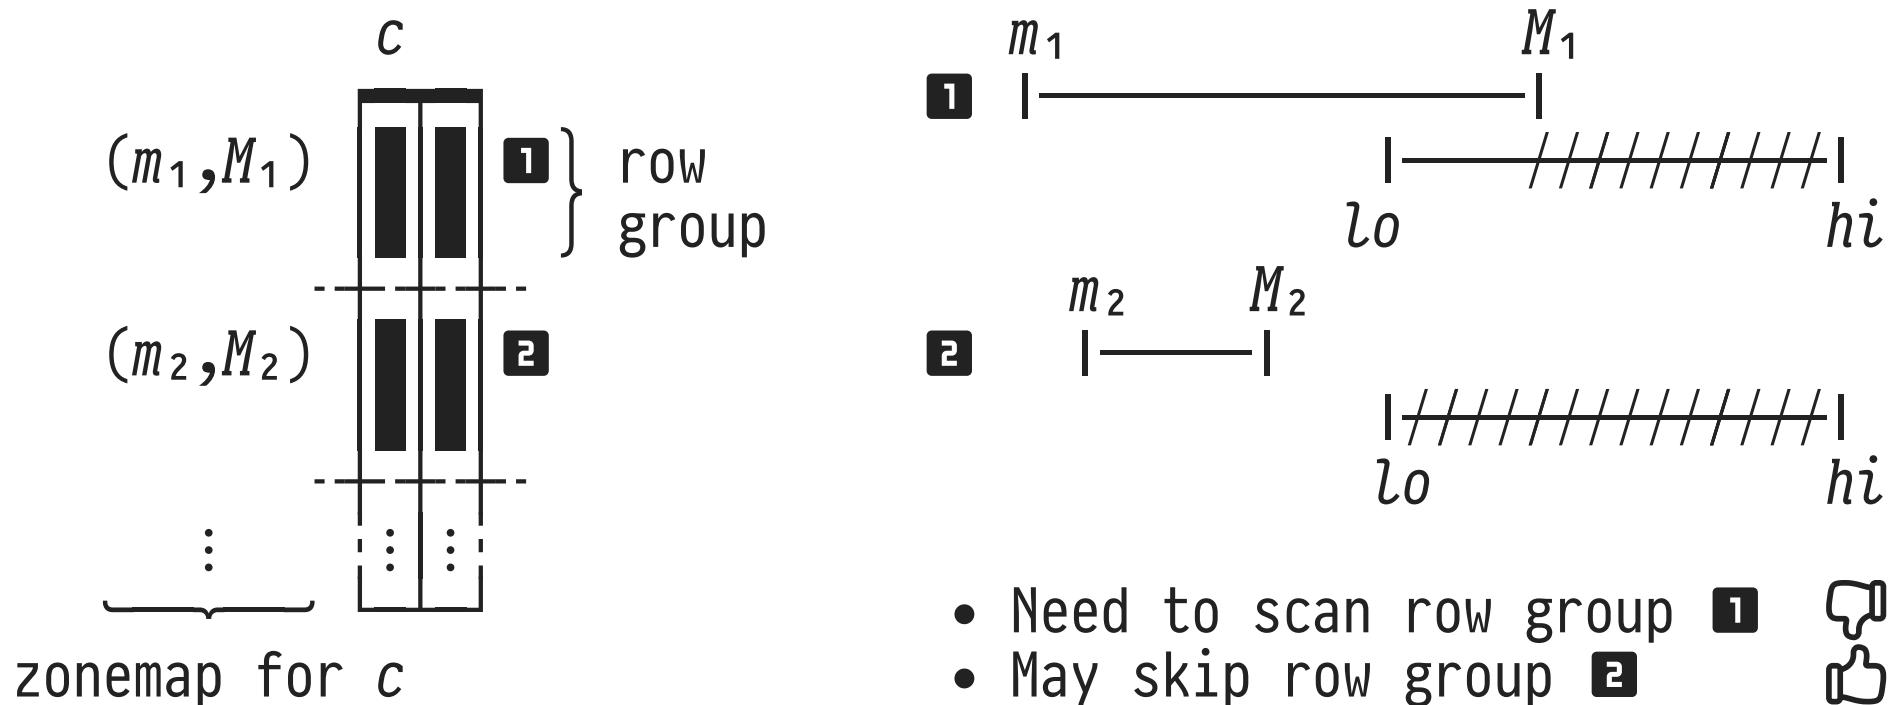
DuckDB implements two kinds of **indexes**:

1. **Zonemaps** (also: *min-max indexes*) are integral part of the columnar table storage .
 - Automatically created, always present for all columns.
 - Used when predicates are pushed down into **Sequential Scan**.
2. **Adaptive radix trees (ART)** are tree-shaped data structures maintained outside/in addition to table storage  + .
- Consume working memory space.
- Require maintenance on table updates.
- Created either
 - manually via SQL's DDL statement **CREATE INDEX** or
 - automatically when constraints **UNIQUE**, **PRIMARY KEY**, and **FOREIGN KEY** are declared for a table.

2 | Zonemaps

Zonemaps are baked into DuckDB's columnar table storage format:

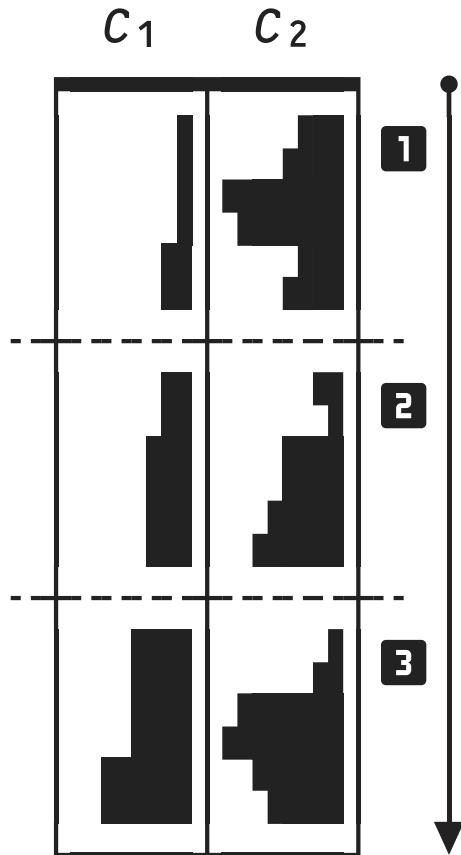
- Each column c is divided into *row groups* of 120K (122880) rows.
- Each row group holds a (\min, \max) entry that—quite roughly—characterizes the contained values.
- Operator **Sequential Scan** can safely **skip row groups** for which a predicate $lo \leq c \leq hi$ will always fail (:///):



Zonemaps and Column Ordering

Column ordering affects the effectiveness of zonemaps.

01 #015



- c_1 :** In each row group, the zonemap entries (\min, \max) for c_1 have a small span $l—l$.
 \Rightarrow A large number of row groups will never satisfy $lo \leq c_1 \leq hi$. Can skip.
- c_2 :** All spans in unordered column c_2 are wide $l—————l$ and cover the active domain of c_2 .
 \Rightarrow Most (all) row groups will potentially contain hits for $lo \leq c_2 \leq hi$. No skipping.

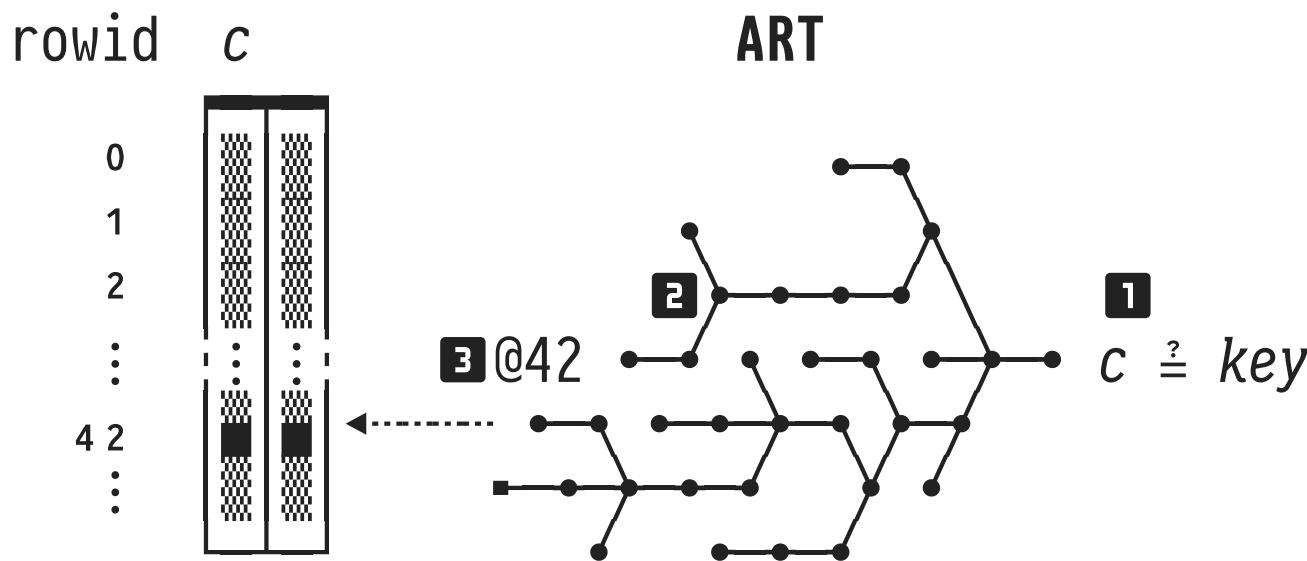
⌚ Keeping base tables ordered may help predicate evaluation.

3 | Adaptive Radix Tree (ART) Indexes

ART indexes are **ordered search tree structures** built to evaluate simple **equality and range predicates** (assume an ART on column c):

$c = \text{key}$ | $c \text{ IN } (\text{key}_1, \text{key}_2, \dots)$ | $c < \text{key}$ | $c > \text{key}$

- ARTs and tables are separate. These extra data structures
 - use **row IDs** to refer to rows in their associated table,
 - must be created, then maintained under table updates, and
 - occupy space in working memory.



Index lookup:

-
- 1 Enter at root
 - 2 Navigate ART to leaf node →
 - 3 Use rowid @42 to access row

Creating ART Indexes in DuckDB

1. Manual:

```
CREATE [UNIQUE] INDEX index ON table (column [, column, ...]);
DROP INDEX [IF EXISTS] index;
```

- **NB.** More indexes aid query evaluation but incur maintenance and space¹ overhead. A tradeoff in physical database design.
2. **Implicit** (supports constraint enforcement). DDL statement **1** creates on a unique ART index on *t(c)* behind the scenes:

```
1 CREATE TABLE t (... , c ... PRIMARY KEY, ...); -- also: UNIQUE
2 CREATE TABLE s (... , f ... REFERENCES t(c), ...);
```

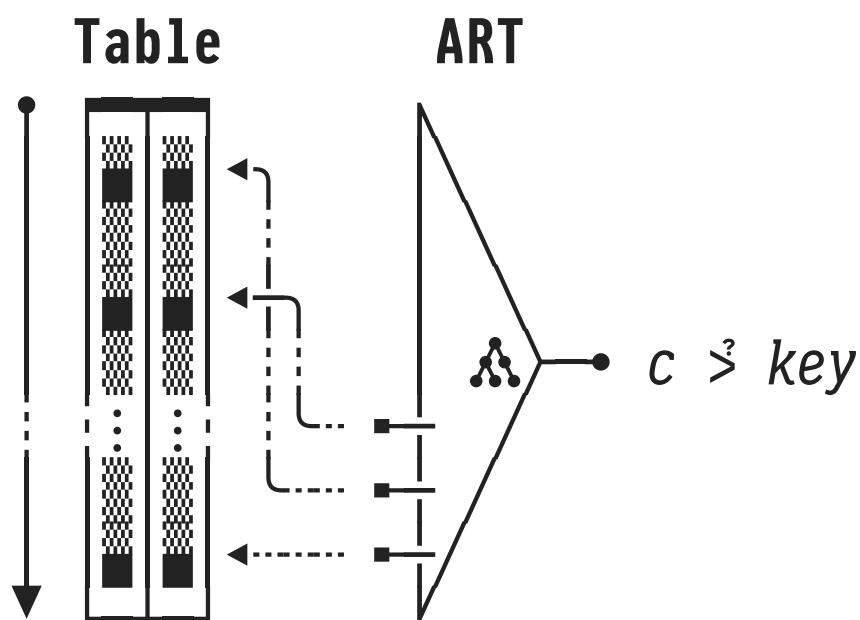
- DDL statements like **INSERT INTO s VALUES (... , key, ...)** or **UPDATE t SET c = key** lead to lookups *c* \doteq *key* in that index.

¹ As of DuckDB 1.4, memory occupied by ART indexes is under control of the buffer manager but cannot be evicted. Work to rectify this is underway as I type this.

Use the Index? Ignore the Index?

An index lookup predicate p may yield *multiple rows* in table t :

- $p \equiv c \stackrel{?}{=} \text{key}$: If c is not unique in t : a single leaf node \rightarrow holds multiple rowids $\rightarrow \rightarrow \rightarrow$.
- $p \equiv c > \text{key}$: 1 perform lookup $c \stackrel{?}{=} \text{key}$, 2 access rows using the rowids in all leaves $\dashrightarrow \dashrightarrow \dashrightarrow$ following the found leaf:



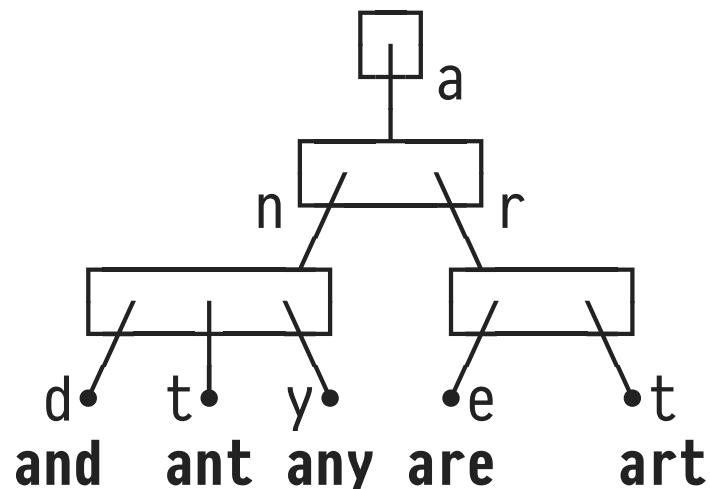
- Lookup finds one leaf (or adjacent leaves $\dashrightarrow \dashrightarrow \dashrightarrow$). #016
 - Yet the rowids $\rightarrow \rightarrow$ may point all over table t . May need to “jump around” to collect rows. Violates memory locality. 👎
- 💡 Use index only if **selectivity** $\text{sel}(p)$ of predicate p is low:

$$0 \leq \text{sel}(p) \stackrel{\text{def}}{=} (\text{SELECT count(*) FILTER (p) / count(*) FROM } t) \leq 1$$

4 | The Internals of Adaptive Radix Trees

An ART on $t(c)$ uses the **bit-wise representation of values** in column c to organize itself. Values are *not* hashed or compared.

- Divide values into groups of s bits each (s is the *span*).
- During tree traversal, the next s bits of *key* determine the child node to descend into.



- Example: $c :: \text{text}$, span $s = 8$ bits.
- Values have $k = 24$ bits.

s bits	$\text{ant} < \text{any}$
$\text{ant} =_2 01000001 \ 01001110 \ 01001110$	$\overbrace{^{\text{common prefix}}}$
$\text{any} =_2 01000001 \ 01001110 \ 01011001$	$\overbrace{^{\text{common prefix}}}$

common prefix \equiv shared path in ART

ART Internals: Mapping Values to Bit Sequences

Map values to bit sequences whose **lexicographic order** properly reflects value sort order:

- **Unsigned ints**: use standard binary representation (little-endian machines: reverse byte order so that MSB comes first).
- **Signed ints** (in two's complement): flip the sign bit, then treat like unsigned ints.
- **IEEE 754 floats**: 1 always flip the sign bit, 2 if the sign bit was originally set, now flip all bits.
- **Text strings**: map UTF-8 byte sequence to binary², end strings with \u2044 (values must not be prefixes of other values).
- **Composite values** (v_1, \dots, v_n): map the individual fields v_i , then concatenate the resulting bit sequences.
- **NULL**: increase bit length k (e.g., by one bit/one byte) to accomodate the additional value.

01 #017

² DuckDB uses function `ucol_getSortKey()` of the C/C++ library `ICU` to perform this mapping.

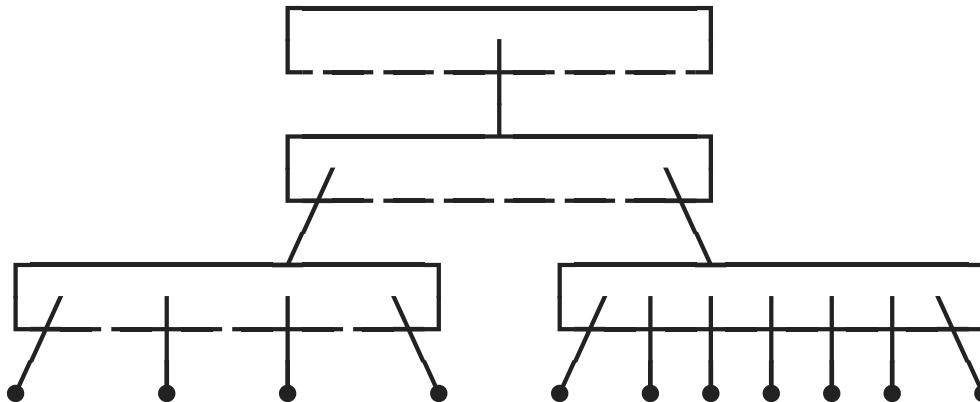
ART Internals: What is a Good Span?

- The height of ARTs depend on the value bit length k (not the number n of entries).
 - An ART for k -bit values has $\lceil k/s \rceil$ levels of inner nodes.
 - ART lookup and insert operations have complexity $O(k)$.
 - ARTs vs. binary search trees (BSTs):
 - Height: if $n > 2^{k/s}$, ARTs have smaller height than BSTs ($\log_2(n)$).
 - Lookup: for values of k bits (k large), comparison ($<$) is $O(k)$. Complexity of BST lookup thus is $O(k \cdot \log_2(n))$.
- 💡 Optimize ART **performance**: Work with a large span s ! 
- ⚠️ But how about the **space usage** of such ARTs? 

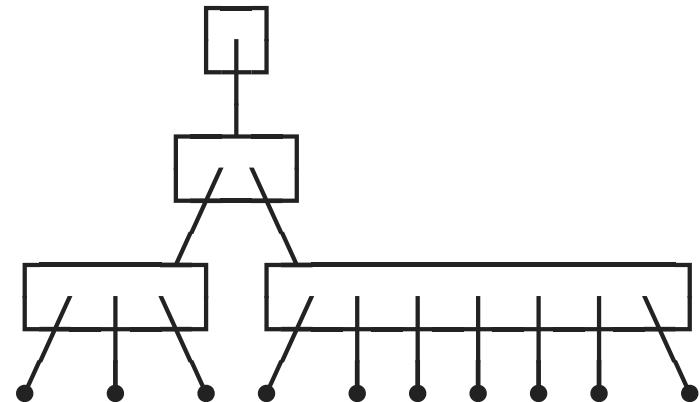
ART Internals: Designing Inner Nodes

- A simple representation of inner ART nodes:
 - Array of 2^s child pointers. (Nodes would have **fixed size**.)
 - Use s -bit chunk of *key* to index array + access child node.
 - But: Most child pointers will be NULL (—), space usage will be excessive if s is large (grows exponentially):

Node size fixed



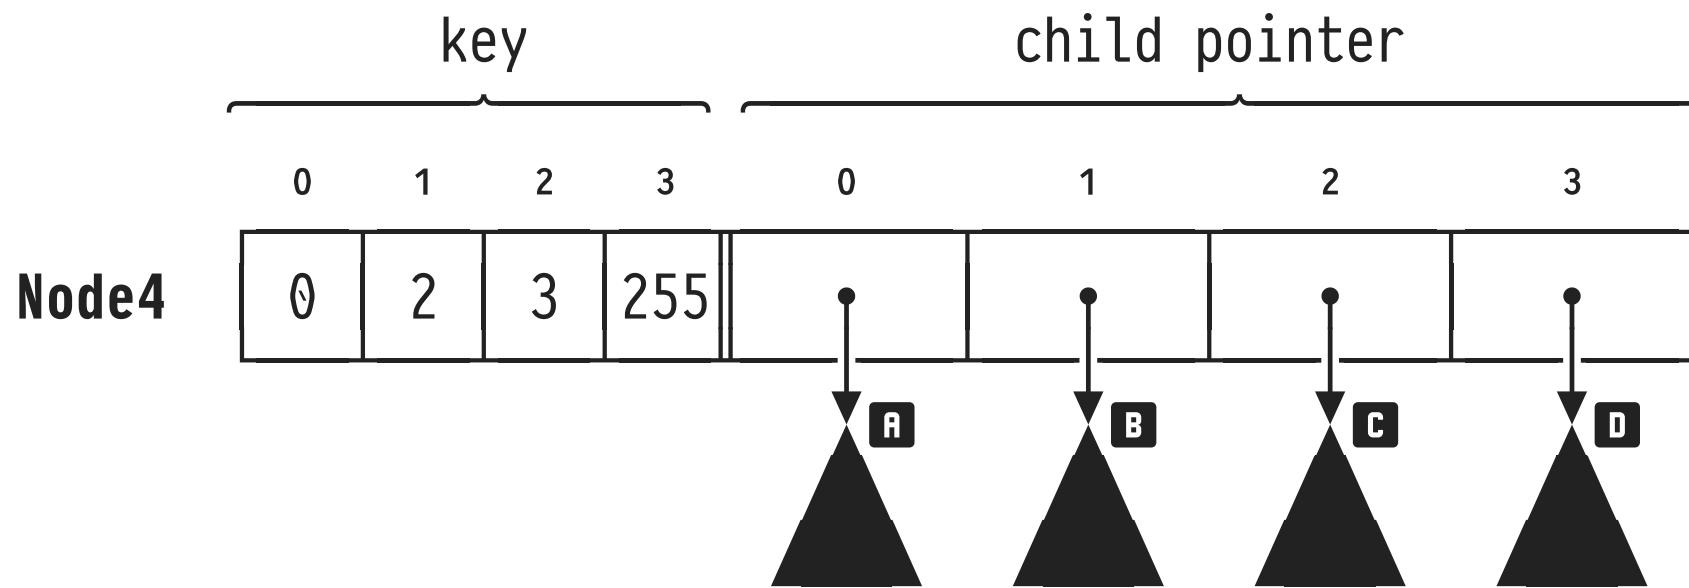
Node size adaptive



⌚ Use a fixed large s but **adapt node size** (use variable fanout).

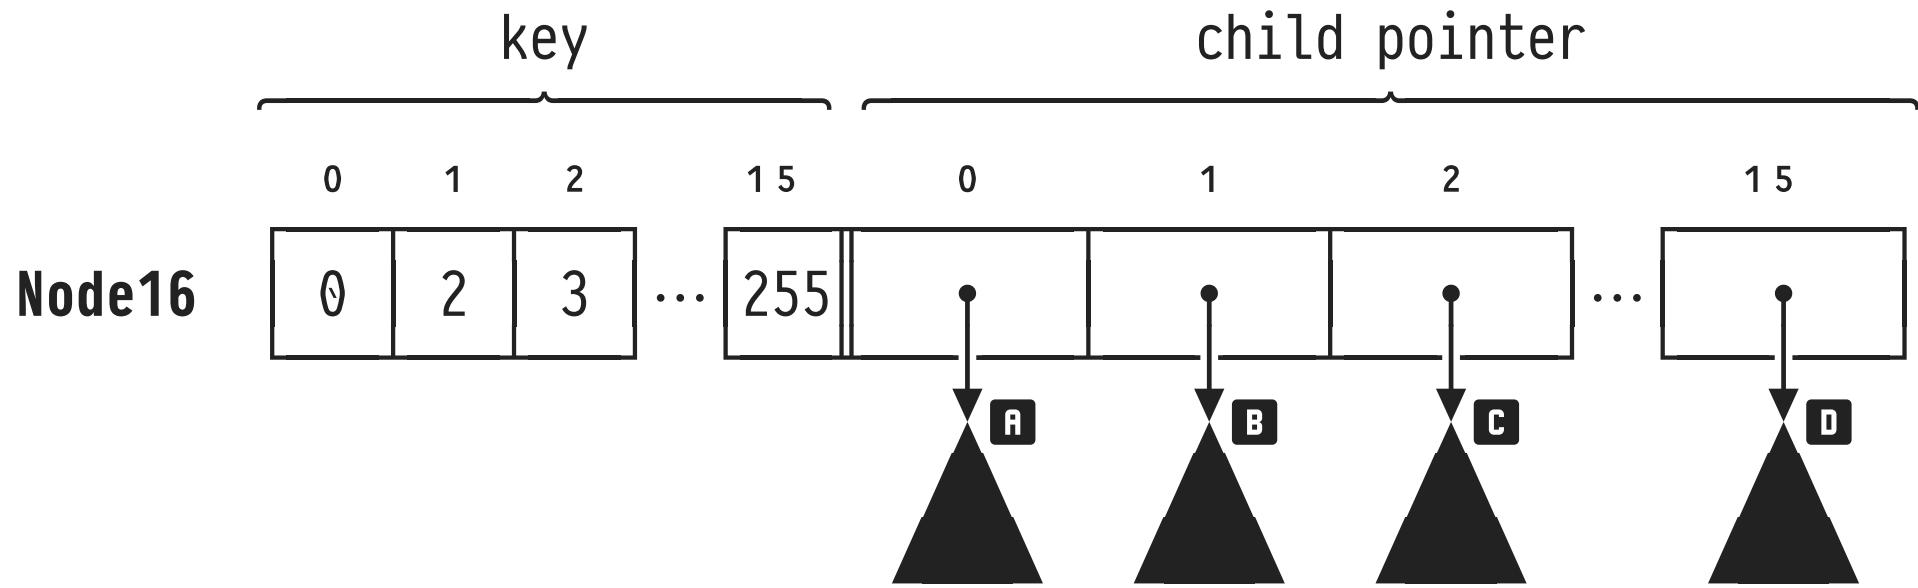
ART Internals: Four Inner Node Types 1

- DuckDB: $s = 8$ (can cut values into bytes, admits large fanout).
 - Choose inner node type based on the number of non-NULL child pointers: 4, 16, 48, or 256 ($= 2^s$).
 - On value insertion/deletion, switch to larger/smaller node type when node overflows/underflows.



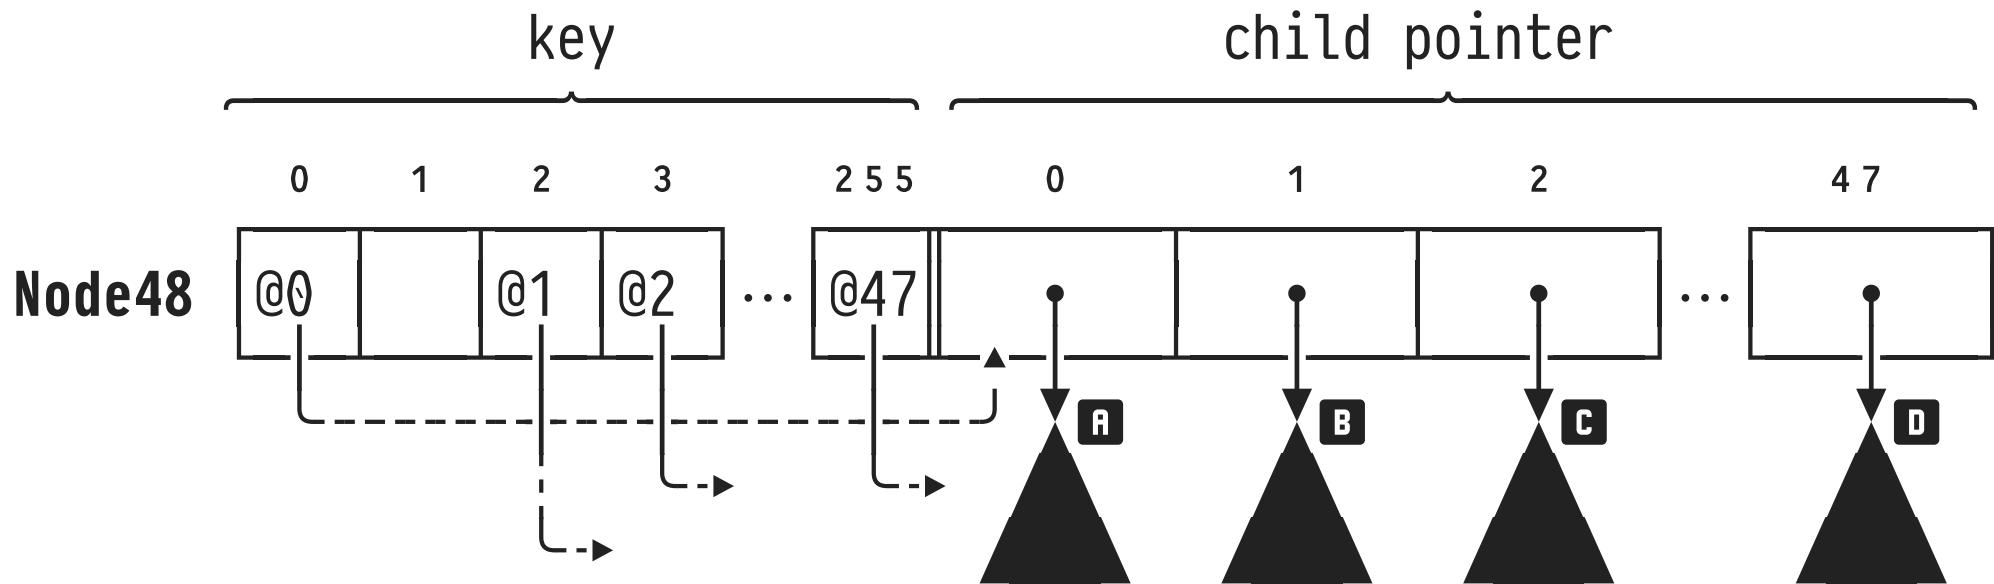
- **Node4:** Array of 4 keys is ordered (search left to right).
 - Value/pointer pairs stored at corresponding indices 0...3.

ART Internals: Four Inner Node Types ②



- **Node16:** stores 5...16 child/pointer pairs.
- Locate current *key* byte by binary search or parallel SIMD comparisons.

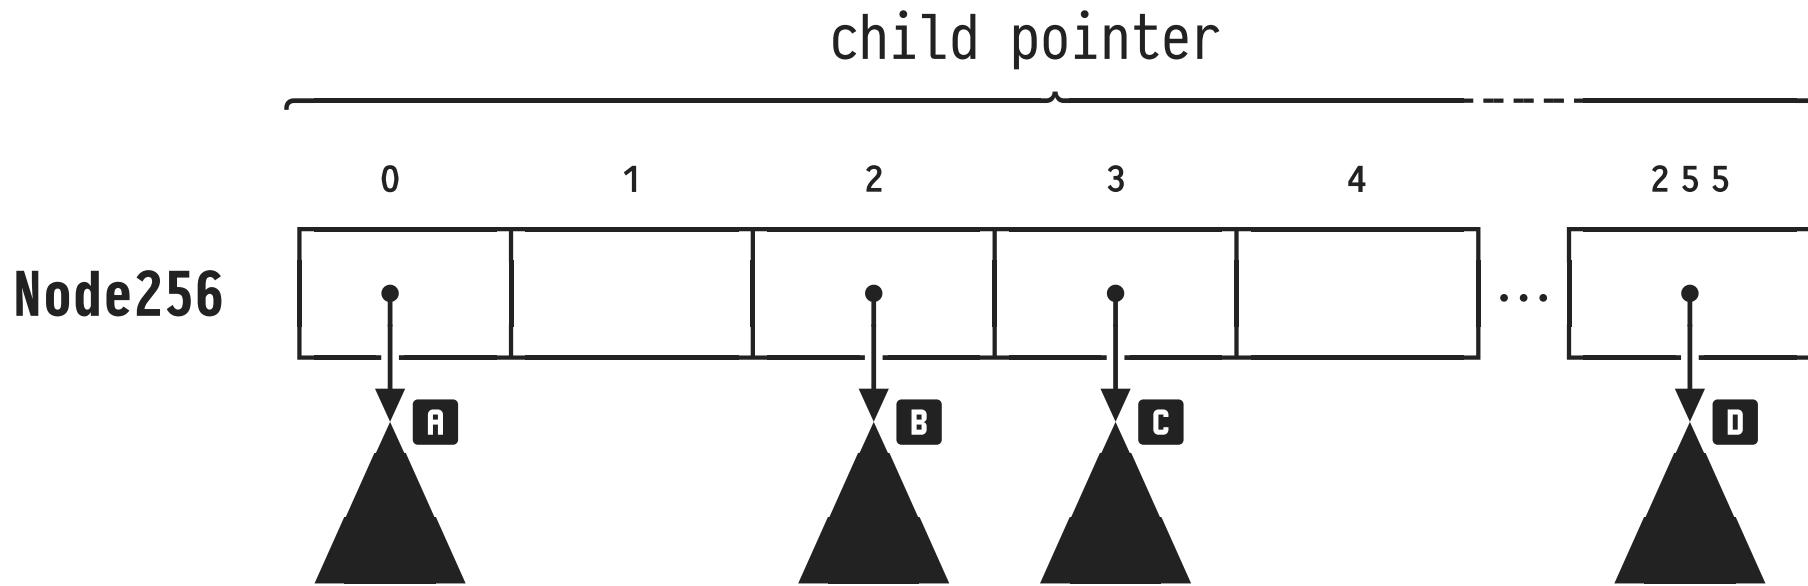
ART Internals: Four Inner Node Types ③



- **Node48:** With 17...48 values in a node, searching the value array becomes expensive. Thus: do *not* store values. Instead:
 - Use current byte of *key* as index into array of 256 indices.
 - *key[<key byte>]* holds the index *@i* $\in \{0 \dots 47\}$ of the corresponding child pointer. Access *pointer[@i]*.
- Saves space compared to an array of 256 8-byte pointers:

$$640\text{B} = 256\text{B} + 48 \times 8\text{B} < 256 \times 8\text{B} = 2048\text{B}.$$

ART Internals: Four Inner Node Types ③

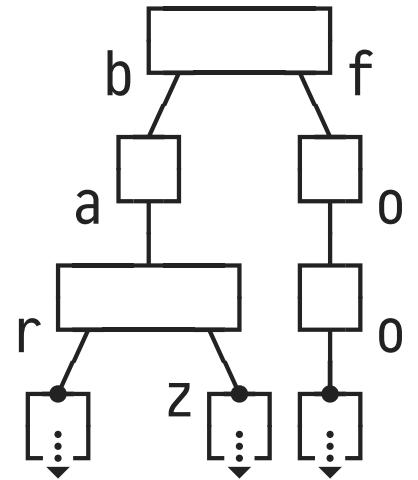


- **Node256:** If an inner node needs to hold 49...256 entries, invest $256 \times 8B = 2kB$ to hold an array of 256 child pointers.
 - Simply return `pointer[<key byte>]`.

ART Internals: Leaf Nodes

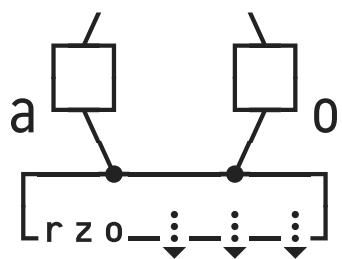
In an ART for $t(c)$, **leaf nodes** hold rowids that point to the rows in table t in which column c holds the search value key .

- Possible ART leaf designs:



(This is an ART for values **bar**, **baz**, **foo**.)

} **Single-entry leaves**, each hold rowids (↓) for one value val .



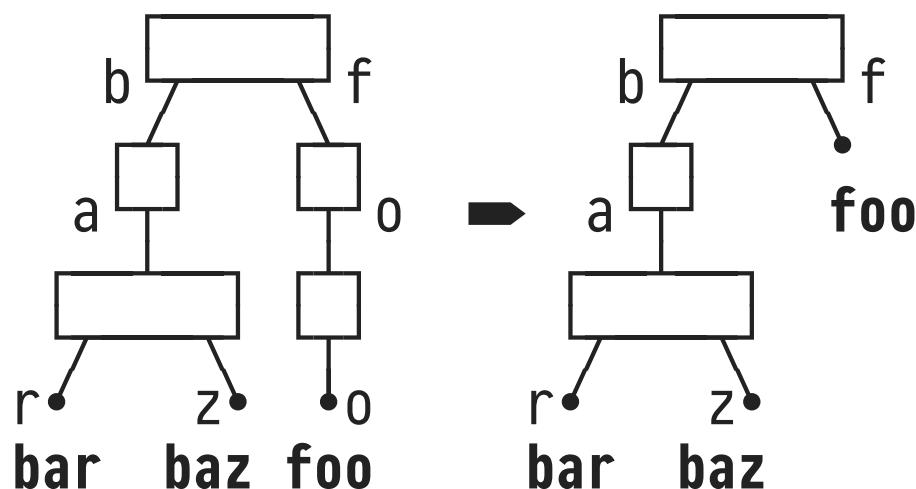
} **Multi-entry leaf**, holds rowids for multiple values. DuckDB has Leaf7/Leaf15/Leaf256 much like Node4...256 (hold rowids, not pointers).

ART Optimizations: Lazy Expansion + Path Compression

To efficiently use space and save tree traversal effort during lookups, try to remove inner ART nodes—and thus **reduce tree height**—if possible.

Most effective if values are long (bit length k is large).

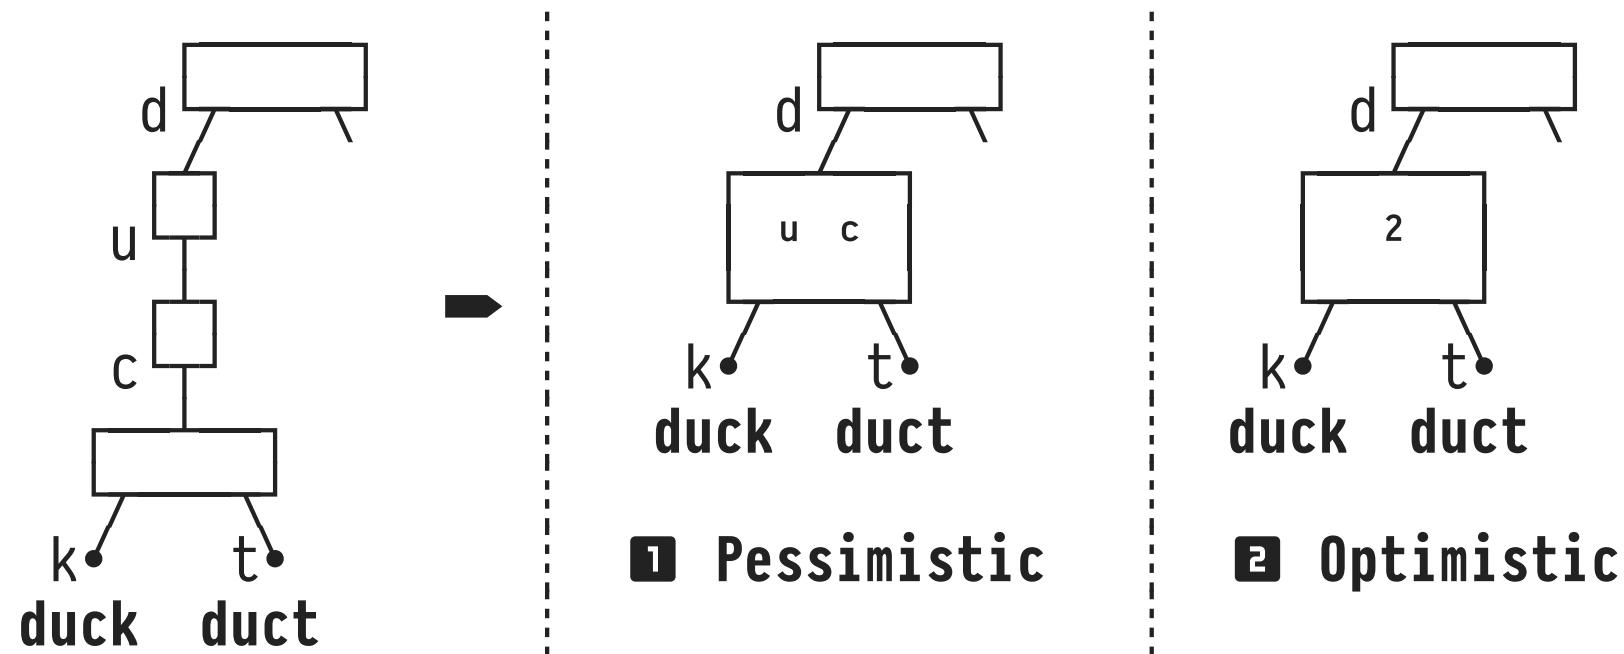
1. **Lazy expansion:** create inner nodes only if they are required to distinguish at least two leaf nodes.



- Saves two inner nodes on the path to leaf **foo**.
- Expand path if another leaf with prefix **f** is added.
- Path does not spell **foo**: store at leaf or in table.

ART Optimizations: Lazy Expansion + Path Compression

2. **Path compression:** remove inner nodes that have a single child.



- ① Inner node holds byte sequence (here: **uc**) of preceding single-child nodes. Compare this sequence against **key**.
- ② Inner holds count (**2**) of removed single-child nodes. Skip that many bytes in **key**. At a leaf, compare its value to **key** to ensure that traversal indeed reached the proper leaf.

ART: Index Lookup (Lazy Expansion + Pessimistic Path Compression)

search(node, key, depth):

```

1 if node = NULL
2   | return NULL
3 if isLeaf(node)
4   | if leafMatches(node, key)
5   | | return node
6   | return NULL
7 s  $\leftarrow$  byteSequence(node)
8 if s  $\neq$  key[depth...depth+ $\text{len}(s)$ ]
9   | return NULL
10 depth  $\leftarrow$  depth +  $\text{len}(s)$ 
11 child  $\leftarrow$  findChild(node, key[depth])
12 return search(child, key, depth+1)

```

search failed (findChild() returned NULL)	1
reached leaf level?	2
reached the proper leaf?	3
yes, success	4
no, failure	5
byte sequence in inner node	6
does partial key match	7
byte sequence?	8
skip to next key byte	9
key byte selects child	10
descend into subtree	11

- Invoke with **search(node: <ART root>, key: 'duck', depth: 0)**.
- Returns leaf node that holds $\text{rowid}(s)$ or NULL.

5 | Making the Most of Indexes

DuckDB is an OLAP DBMS, optimized to scan entire tables. Basic index support is present—but there certainly is unused potential.

- The following predicates *could* be evaluated by Index Scan:³

Predicate	Index Constellation	Index used by  ?
$c_1 \text{ } \textcircled{<} \text{ } val_1$	$t(c_1)$	 (if selective)
	$t(c_1, c_2)$ (composite)	
$c_1 \text{ } \textcircled{<} \text{ } val_1 \text{ AND } c_2 \text{ } \textcircled{<} \text{ } val_2$	$t(c_1, c_2)$ (composite)	
	$t(c_1)$ and/or $t(c_2)$	
$c_1 \text{ } \textcircled{<} \text{ } val_1 \text{ OR } c_2 \text{ } \textcircled{<} \text{ } val_2$	$t(c_1)$ and $t(c_2)$	
$c_1 \text{ } \texttt{LIKE} \text{ } 'p_tt?rn'$	$t(c_1)$	

Index support in DuckDB (as of version 1.4)

- (PostgreSQL  is optimized to operate as an OLTP DBMS and can use indexes in all of these scenarios.)

³ In this table,  represents a SQL comparison operator: $<$, \leq , $=$, \geq , $>$.

Index Traversals

Index traversals (beyond equality-based index lookups for a key val) can support **range predicates**:

01 #019

- $c_1 \geq val_1$ or $c_1 \geq val_1 \text{ AND } c_1 \leq val_2$: perform lookup \rightarrow for val_1 , then scan leaves left to right \dashrightarrow (until $> val_2$). 1
- $c_1 \text{ LIKE } 'prefix\%'$: perform lookup \rightarrow for $prefix$, hit inner node. Scan leaves in subtree \dashrightarrow . 2

